

REFINED ARCHITECTURE SPECIFICATION

Global Earning Pages Architecture Refinement

Resolving six open issues in the platform's pivot from a Christmas-only site to a global creator-publishing and traffic-monetization infrastructure.

This document addresses ad-network co-display rules, GDPR/CCPA consent management, bot detection and traffic-quality scoring, platform-side revenue models, earnings-claim compliance language, and account lifecycle / offboarding. Each section provides the architectural decision, the data model changes required, and the operational controls needed for production.

DOCUMENT TYPE

Technical Architecture Refinement v2

STATUS

For Engineering Review

Table of Contents

Chapter 1 • Executive Summary	3
1.1 What changes from the original concept	3
1.2 What does not change	3
Chapter 2 • Ad-Network Co-Display Rules	5
2.1 Findings from publisher policies	5
2.2 Architectural decision: Policy-Driven Placement Engine	5
2.3 Data model additions	6
2.4 Runtime placement algorithm	6
Chapter 3 • GDPR/CCPA Cookie Consent Architecture	8
3.1 Required CMP integration	8
3.2 Consent flow architecture	9
3.3 Data model additions	9
Chapter 4 • Traffic Quality and Bot Detection	11
4.1 Signal stack	11
4.2 Scoring algorithm	12
4.3 Tools and integration points	12
4.4 Moderation state transitions	12
Chapter 5 • Platform-Side Revenue Model	14
5.1 Revenue layer comparison	14
5.2 Rev-share decision: none, by default	14
5.3 Incentive alignment	15
5.4 Launch sequencing	15
Chapter 6 • Earnings-Claim Compliance Language	16
6.1 Recommended language changes	16
6.2 Required disclaimers	16
6.3 Analytics display rules	17
Chapter 7 • Account Lifecycle and Offboarding	18
7.1 Ad-integration lifecycle states	18
7.2 Offboarding scenarios	18
7.2.1 User's ad-network account is suspended	19

7.2.2 User deletes their platform account	19
7.2.3 User wants to migrate from one network to another	19
7.2.4 Platform bans a user	19
7.3 Data retention policy	19
Chapter 8 • Consolidated Data Model Changes	21
8.1 Issue 1 — Co-display rules	21
8.2 Issue 2 — Consent layer	21
8.3 Issue 3 — Trust Score	21
8.4 Issue 4 — Revenue model	21
8.5 Issue 5 — Compliance language	21
8.6 Issue 6 — Lifecycle & offboarding	21
Chapter 9 • Updated Rollout Sequence	23
9.1 Phase 1 — Minimum viable platform (MVP)	23
9.2 Phase 2 — Multi-network and moderation automation	23
9.3 Phase 3 — Direct sales and discovery	23
Chapter 10 • Remaining Open Items	25
10.1 Ad-network API integration depth	25
10.2 Additional ad networks	25
10.3 Mobile app vs mobile web	25
10.4 Tax and reporting obligations for subscriptions	25
10.5 Content licensing and IP	25

Chapter 1 • Executive Summary

This document refines the Global Earning Pages architecture by resolving six open issues raised during the original concept review. Each issue is addressed with a concrete architectural decision, the data model changes required to support it, and the operational controls that must accompany it. The objective is to convert the original concept from a defensible idea into an implementable engineering specification that an engineering team can build against without further clarification rounds.

The six issues are: (1) ad-network co-display rules, where Adsterra and Monetag impose restrictions on competitive ads, ad density, and content adjacency that the placement engine must enforce; (2) GDPR/CCPA cookie consent, where the platform becomes jointly liable for ad-tag cookies set on user pages and must integrate an IAB TCF v2.2-certified Consent Management Platform; (3) traffic-quality scoring, where the proposed Trust Score must be backed by real signals rather than aesthetic placeholders; (4) the platform-side revenue model, where the original concept left the buyer-side economics undefined; (5) earnings-claim compliance language, where the "Earn Money" framing creates implied-earnings exposure under FTC and analogous regimes; and (6) account lifecycle and offboarding, where the system must handle banned ad-network accounts, GDPR right-to-erasure requests, and platform-initiated suspensions without leaving orphaned ad code on live pages.

The unifying architectural principle is strict separation between the User Ad Engine and the Platform Ad Engine, with the Consent Layer as the single gate through which both must pass. This separation is what makes the business model defensible: the platform never pays users, never processes ad payouts, and never becomes the regulatory counterparty for the user's ad-network relationship. The platform's role is infrastructure, moderation, and discovery — not financial intermediation.

1.1 What changes from the original concept

The original concept proposed a "GLOBAL EARNING PAGES" platform with an "Earn Money" dashboard button, an ad-placement engine that allowed user ads and platform ads to coexist on the same page without constraint, a Trust Score field with no defined inputs, and no specified offboarding flow. Each of these is addressed in the chapters that follow, but the headline changes are: the dashboard button is renamed "Monetization"; the ad-placement engine becomes a policy-driven component that consults per-network compatibility rules before rendering; the Trust Score is decomposed into five sub-scores each backed by specific signal sources; and every ad integration is registered with a lifecycle state machine that handles disabling, revocation, and deletion as first-class operations.

1.2 What does not change

The core product thesis remains intact. The platform is still a global creator publishing and traffic-monetization infrastructure. Users still create Special Pages, publish content, share URLs, and optionally connect their own approved ad networks. Christmas 2026 remains the flagship launch campaign rather than the product identity. The page builder, content blocks, sharing engine, creator levels, traffic challenges, documentation center, and admin moderation tools described in the original concept are preserved without material change. The refinements here tighten the compliance and reliability of the monetization layer without altering the user-facing product surface area.

Architectural invariant: User Ads and Platform Ads share page real estate but never share consent state, network account, or revenue flow. The Consent Layer is the only component permitted to gate either engine.

Chapter 2 · Ad-Network Co-Display Rules

The original concept assumed that user-supplied Adsterra or Monetag tags and platform-supplied ad inventory could be interleaved on the same Special Page without conflict. Web research and the published policies of both networks indicate that this assumption is partially correct but requires enforcement. Adsterra explicitly permits its ads to be displayed alongside Google AdSense and other ad networks on the same page, but publishers must avoid "excessive ads" and ensure placements do not conflict. Monetag similarly supports multiple ad formats simultaneously on a single site and allows other networks alongside, but imposes its own content and traffic rules that the platform cannot override.

The implication for the placement engine is that co-display is generally permitted but is not unconditional. The platform must enforce three classes of constraint: (a) ad-density caps to avoid the "excessive ads" trigger that Adsterra and most major ad networks prohibit; (b) per-network compatibility rules, since individual networks may prohibit specific competitive networks or specific creative formats adjacent to their own; and (c) content-adjacency rules, where ads must not be placed next to content categories that violate either network's publisher content policies.

2.1 Findings from publisher policies

Adsterra's publisher documentation and several third-party analyses confirm that Adsterra ads can be shown alongside AdSense and other networks, but the publisher must avoid excessive ad placements and ensure placements do not conflict with AdSense policies when co-displayed. Monetag's documentation describes support for multiple simultaneous ad formats on a single site and does not impose a general exclusivity rule against other networks. However, both networks reserve the right to suspend publishers for low-quality traffic, prohibited content categories, or violation of their respective terms — and these suspensions are against the user's publisher account, not the platform.

Constraint	Source	Platform enforcement
No "excessive ads" on a single page	Adsterra policy; general industry norm	Cap total ad units per Special Page (default: 4); reject placements beyond cap
Per-network competitive exclusions	Per-network publisher ToS (varies)	Maintain AdNetworkCompatibility table; placement engine consults it before rendering
Content-adjacency restrictions	Both networks prohibit certain content categories	Page moderation state gates ad rendering; Restricted/Suspended pages serve no ads
Traffic-quality requirements	Adsterra: no bots/incentivized; Monetag: no spoofing/fraud	Traffic-Quality sub-score below threshold disables ad rendering on the page
Site-verification status	Monetag requires website verification before serving zones	User onboarding flow requires site verification step before integration approval

Table 2.1 — Co-display constraints and platform enforcement points

2.2 Architectural decision: Policy-Driven Placement Engine

The placement engine is redesigned as a policy-driven component rather than a simple slot renderer. When a Special Page is requested, the engine collects the set of intended ad placements (both user-supplied and platform-supplied), then consults a policy matrix to determine which placements may render together, in what order, and with what density. The policy matrix is admin-configurable so that changes in network rules can be propagated without code changes.

The policy matrix is stored as an `AdNetworkCompatibility` table mapping pairs of ad networks (including "Platform" as a synthetic network) to a compatibility verdict: `ALLOWED`, `ALLOWED_WITH_LIMITS`, or `FORBIDDEN`. When the engine encounters an `ALLOWED_WITH_LIMITS` verdict, it consults the associated limit record for maximum simultaneous units, minimum vertical separation, and required content class between the two placements. `FORBIDDEN` verdicts cause one of the two placements to be suppressed based on a priority order specified at the placement level.

2.3 Data model additions

```
AdNetwork
id Int PK
code String unique // "adsterra", "monetag", "platform"
displayName String
integrationTypes String[] // SCRIPT, DIRECT_LINK, NATIVE, BANNER, ...
requiresSiteVerification Bool
policyDocUrl String?
status AdNetworkStatus // ACTIVE, DEPRECATED, BANNED

AdNetworkCompatibility
id Int PK
networkA AdNetwork FK
networkB AdNetwork FK
verdict CompatibilityVerdict // ALLOWED, ALLOWED_WITH_LIMITS, FORBIDDEN
maxSimultaneousUnits Int?
minVerticalSeparationPx Int?
requiredContentClassBetween String? // e.g. "user_content_min_200px"
notes String?
updatedAt DateTime

AdPlacementPolicy // global platform-level caps, applied after pairwise checks
id Int PK
maxAdUnitsPerPage Int // default 4
maxPlatformAdsPerPage Int // default 2
maxUserAdsPerPage Int // default 2
minContentBetweenAdsPx Int // default 200
updatedAt DateTime
```

2.4 Runtime placement algorithm

When a Special Page is rendered, the placement engine executes the following sequence. First, it loads the page's intended placements sorted by priority. Second, it walks the list and for each placement, checks pairwise compatibility with every placement already accepted; if any pair is `FORBIDDEN`, the lower-priority placement is dropped. Third, it applies the global caps from `AdPlacementPolicy` (total units, platform units, user units, minimum content separation). Fourth, it renders the surviving placements in priority order, injecting spacer content blocks between them where required by the `requiredContentClassBetween` rule.

Operational note: the compatibility table is the single point of truth for co-display rules. When Adsterra or Monetag updates their publisher ToS, the admin updates one row — no code deploy required. The placement engine reads the table on every render (cached with a 60-second TTL) so policy changes propagate within one minute.

Chapter 3 • GDPR/CCPA Cookie Consent Architecture

When users embed Adsterra or Monetag tags on their Special Pages, the resulting cookies and tracking pixels are set in the visitor's browser against the platform's domain. Under the General Data Protection Regulation (GDPR), consent must be prior, informed, specific, unambiguous, and revocable — implied consent is unacceptable. The California Consumer Privacy Act (CCPA), as amended by the CPRA, follows an opt-out model rather than an opt-in model but, with CPRA's dark-patterns rules fully in effect from January 2026, requires a functional and symmetric opt-out mechanism. The platform, as the controller of the page chrome and the entity that permits third-party ad tags to be embedded, is jointly liable for lawful consent collection.

The industry-standard mechanism for signaling consent to ad networks is the Interactive Advertising Bureau's Transparency and Consent Framework version 2.2 (IAB TCF v2.2). Google requires publishers serving ads in the EU/UK to use a Google-certified CMP that integrates with TCF v2.2. Adsterra and Monetag both document compatibility with the TCF signal, meaning a correctly implemented CMP will suppress their tags until consent is granted, without requiring per-tag patching by the platform.

3.1 Required CMP integration

The platform must integrate a CMP certified under IAB TCF v2.2. The CMP is embedded once in the Special Page chrome and exposes a JavaScript API (`__tcfapi`) that ad tags consult before executing. Certified CMPs that may be considered include Usercentrics, CookieYes, Seers, Pandectes, TrustArc, and Setupad; the choice depends on pricing model, regional coverage, and the specific TCF v2.2 features supported. The integration is platform-wide and not configurable per user — every Special Page renders the same CMP, with the same consent strings, signaling the same set of vendor consent purposes.

CMP requirement	Why it matters	Implementation point
IAB TCF v2.2 certified	Required by Google for EU/UK ad serving; recognized by Adsterra/Monetag	Select only from IAB-certified list; verify annually
Google-certified (if Google Ads used)	Required for Google Ads serving in EU/UK	Cross-check Google's certified CMP list at integration time
Per-vendor granular consent	Visitors must be able to consent per purpose and per vendor	CMP configured with all active ad-network vendors registered
Symmetric opt-out (no dark patterns)	CPRA 2026 enforcement; "Don't Sell" must be as easy as "Accept"	Equal-weight buttons; no pre-ticked boxes; no "Continue" pressure
Consent string persistence	Subsequent page loads must honor prior consent state	CMP stores TC string in first-party cookie; platform reads on SSR
Server-side signal forwarding	Some networks require server-to-server consent signals	Platform forwards TC string in ad-tag request headers

Table 3.1 — CMP integration requirements

3.2 Consent flow architecture

The consent flow is implemented as a distinct architectural layer that sits between the page chrome and both ad engines. On initial page load, the server renders the Special Page with ad slots marked but no ad tags emitted. The CMP script loads synchronously in the document head, presents the consent banner if no prior consent string is found, and exposes `__tcfapi`. Once the visitor completes the consent flow (or if a valid prior consent string is detected), the CMP fires the consent-updated event; the platform's Ad Renderer listens for this event and emits the appropriate ad tags for each slot, with the TC string forwarded in the request.

For visitors outside GDPR/CCPA jurisdictions, the CMP still loads but operates in a default-grant mode governed by the visitor's IP-derived jurisdiction. This is a deliberate choice: it is cheaper to render the CMP everywhere than to maintain a separate code path that misses edge cases (VPN users, geoip database drift, new state privacy laws). The default-grant mode is logged so that auditors can verify the platform did not silently suppress consent collection for users who should have received it.

```
// Pseudocode: Special Page SSR consent flow
function renderSpecialPage(page, visitor) {
  const consentState = cmp.loadConsentString(visitor.id);
  const jurisdiction = geoip.derive(visitor.ip);
  const requiresConsent = ["EU", "UK", "CA", "VA", "CO", "CT"]
    .includes(jurisdiction);

  return html`
    <head>
      <script src="cmp-loader.js" data-cmp-id="..."></script>
    </head>
    <body>
      ${renderPageContent(page)}
      ${AdSlot placements=[...page.placements]
        gate=consentState
        requiresConsent=requiresConsent
        strategy="defer-until-consent"}
    </body>
  `;
}
```

3.3 Data model additions

```
ConsentRecord // immutable append-only log of consent state per visitor
id Int PK
visitorId String indexed // first-party anonymous ID
tcString Text // IAB TCF v2.2 consent string
jurisdiction String // EU, UK, CA, OTHER
capturedAt DateTime indexed
expiresAt DateTime // 13 months per TCF v2.2
source String // "banner", "prior", "default_grant"

AdNetworkVendorMapping // maps platform ad networks to IAB TCF vendor IDs
id Int PK
adNetworkCode String FK AdNetwork.code
tcfVendorId Int // IAB GVL vendor ID
requiresConsentFor String[] // "advertising", "analytics"
updatedAt DateTime
```

Compliance note: the ConsentRecord table must be append-only. Updates create a new row with a later capturedAt timestamp; the most recent non-expired row governs the current consent state. This supports the GDPR right to demonstrate consent (Art. 7.1) and auditability of consent withdrawals.

Chapter 4 · Traffic Quality and Bot Detection

The original concept proposed a "Trust Score" composed of Risk, Traffic Quality, Content Risk, Ad Risk, and Spam Risk sub-fields, but did not specify the signals feeding each. A score without defined inputs is theater: it cannot defend an account suspension decision, it cannot be audited, and it cannot be improved. This chapter specifies the signal stack that backs each sub-score, the tools used to acquire those signals, and the scoring algorithm that combines them.

Both Adsterra and Monetag explicitly prohibit bot traffic, incentivized traffic, spoofing, fraudulent traffic, and certain misleading or money-making content categories. Adsterra's security guidance identifies bot and incentivized traffic as prohibited and addresses malicious/fraudulent traffic patterns. Monetag likewise prohibits spoofing and fraudulent traffic. Because these prohibitions are against the user's publisher account, the platform's incentive to detect bad traffic early is twofold: to protect the platform's relationship with each network (a high rate of banned users reflects on the platform) and to protect users from being banned for activity the platform could have warned them about.

4.1 Signal stack

Modern bot detection combines multiple techniques because no single signal is sufficient. Industry best practice, as documented by Imperva, Engagelab, and others, combines device fingerprinting, behavioral analysis, traffic analysis, and risk scoring. The platform adopts this layered approach, with each signal feeding one or more sub-scores. Signals are collected both at request time (synchronous, in the edge layer) and asynchronously (post-request, in the analytics pipeline).

Sub-score	Signals	Source / tool
Traffic Quality	Bot score from edge; traffic-source diversity; geographic distribution; returning-visitor ratio; page-view-to-session ratio; click-through rate on placements	Cloudflare Bot Management; first-party analytics; referrer validation
Content Risk	Page-content classification; flagged keywords; moderator decisions; user-report volume; appeal outcomes	Internal content classifier; human moderators; user reports
Ad Risk	Click-velocity on ad placements; ad-click-to-pageview ratio; conversion rate anomalies; repeated identical IPs clicking ads; user's ad-network suspension history	In-house click-fraud pipeline; ad-network webhooks
Spam Risk	Account age; URL creation velocity; URL-share velocity across social platforms; duplicate content fingerprint; known-spam-IP overlap	Account analytics; share-engine logs; content fingerprinting
Account Risk (composite)	Weighted sum of the four sub-scores above plus account-standing modifiers (verification status, prior suspensions, payment-network standing)	Computed nightly by the risk-scoring job

Table 4.1 — Trust Score sub-scores and their signal sources

4.2 Scoring algorithm

Each sub-score is a value in [0, 100] where 100 is best. The composite Trust Score is computed as a weighted sum: Traffic Quality 35%, Content Risk 20%, Ad Risk 25%, Spam Risk 15%, Account Risk modifiers 5%. Weights are admin-configurable and may be tuned without code changes. Sub-scores are recomputed nightly in batch for every active Special Page, and the composite Trust Score is written to the page's moderation state record. Critical signals (e.g. an ad-network webhook reporting account suspension) trigger an immediate out-of-band re-score and a moderation state transition.

Each sub-score is decomposed into signal-level data, so when a moderator inspects a page they can see not just "Traffic Quality = 42" but the underlying signals: 31% of traffic flagged as bot-like by edge; click-through rate 4× platform median; 87% of sessions from a single ASN. This decomposition is what makes the score auditable and what gives moderators the evidence needed to defend a suspension decision to the user.

4.3 Tools and integration points

The platform should adopt a layered tool selection rather than relying on a single vendor. At the edge, Cloudflare Bot Management (or an equivalent such as Imperva or Akamai Bot Manager) provides the first-pass bot score on every request, before the application layer is hit. At the application layer, FingerprintJS (open-source or Pro edition) provides client-side device fingerprinting that catches headless browsers and persistent device-spoofing attempts. At the analytics layer, an in-house click-fraud pipeline processes ad-click events in near-real-time and flags suspicious velocity patterns.

Tool selection should not be over-engineered at launch. The MVP layer is Cloudflare Bot Management (which most production deployments already use for WAF) plus FingerprintJS open-source plus an in-house velocity check on ad clicks. This combination catches the majority of low-sophistication fraud patterns at minimal cost. The advanced patterns (click-farms with rotating residential IPs, ML-driven behavioral mimicry) require ML-based detection that should be deferred until traffic volume justifies the investment.

4.4 Moderation state transitions

The Trust Score does not directly disable monetization. Instead, it drives the moderation state machine: a Special Page's state can be Pending, Approved, Restricted, Suspended, or Banned. Score thresholds trigger state transitions, but every transition is logged with the triggering signals, and any transition that affects monetization (Restricted and below) generates a user-facing notification with the reason and the appeal path. The state machine is designed so that automated state changes can be reversed by a human moderator without losing the audit trail.

```
TrustScore (snapshot per page, written nightly)
id Int PK
pageId Page FK
trafficQuality Int // 0-100
contentRisk Int
adRisk Int
spamRisk Int
composite Int // weighted sum
signalBreakdown Json // per-signal values for audit
computedAt DateTime indexed

ModerationEvent // append-only state-transition log
```

```
id Int PK
pageId Page FK
fromState ModerationState
toState ModerationState
reason String
triggerSignals Json // copy of signals at transition time
triggeredBy String // "automated" or userId
createdAt DateTime indexed
```

Chapter 5 • Platform-Side Revenue Model

The original concept left the platform-side revenue model undefined beyond "platform ads run on eligible Special Pages." This chapter specifies four distinct revenue layers, the buyer-side economics of each, the decision matrix for revenue-share with users, and the incentive-alignment principles that prevent the revenue model from polluting the user-facing product.

The four revenue layers are: (1) direct platform ad inventory sold by the platform itself to advertisers, displayed on eligible Special Pages; (2) platform-supplied ad-network inventory (e.g. the platform running its own Adsterra or Monetag publisher account on platform-managed placements), distinct from user-supplied ad-network integrations; (3) sponsored placements where creators opt to be promoted in discovery surfaces; and (4) creator-tooling subscriptions for power features (advanced analytics, custom domains, premium themes, removal of platform ads on their own pages). Each layer has different economics, different compliance profiles, and different relationships to the user ad layer.

5.1 Revenue layer comparison

Layer	Buyer	Revenue type	User rev-share?	Conflict with user ads?
Direct platform inventory	Advertisers buying from platform sales team or self-serve portal	CPM / CPC / fixed placement	No (user ads run independently)	No — co-exists via placement engine
Platform ad-network inventory	Adsterra / Monetag paying the platform's own publisher account	CPM / revenue share from network	No	No — but capped by AdPlacementPolicy
Sponsored discovery placements	Creators paying for visibility in search/trending/explore	Fixed fee per period; or bid-based	N/A (the creator is the buyer)	No
Creator tooling subscriptions	Creators paying monthly for premium features	Recurring SaaS	N/A	Optional "remove platform ads" tier — only affects platform layer

Table 5.1 — Four platform-side revenue layers

5.2 Rev-share decision: none, by default

The original concept did not specify whether users receive a share of platform-ad revenue generated on their pages. The recommendation here is: no rev-share by default. Rationale: rev-share introduces payment processing, KYC, tax reporting, chargeback handling, and cross-border payout infrastructure — all of which are precisely the burdens the architecture was designed to avoid by making ad networks the party that pays users. Rev-share would convert the platform from infrastructure provider to financial intermediary, with the compliance surface that entails.

The alternative mechanism for incentivizing users to permit platform ads on their page is a creator-tooling reward: users who enable platform ads on their page receive access to premium features at no cost (advanced analytics, custom themes, removal of platform watermark, higher content-block limits). This is a benefit-in-kind exchange rather than a revenue share, and it avoids the regulatory surface of monetary payouts while still giving users a reason to opt in.

5.3 Incentive alignment

Each revenue layer must be checked against three alignment tests before launch. First, does the layer create an incentive for the platform to act against the user's interest? Direct platform inventory sold to advertisers creates pressure to keep ad density high, which conflicts with user experience — this is mitigated by the AdPlacementPolicy caps. Second, does the layer create an incentive for users to act against the platform's interest? Creator tooling subscriptions create pressure to let paying users violate content policies — this is mitigated by making subscription status orthogonal to moderation state. Third, does the layer create an incentive for the platform to discriminate between users in ways that damage discovery? Sponsored discovery placements create pressure to demote non-paying creators in search ranking — this is mitigated by keeping sponsored slots visually and structurally separate from organic discovery.

5.4 Launch sequencing

Not all four layers should launch at the same time. The recommended sequence is: Layer 2 (platform ad-network inventory) at launch, because it requires no sales infrastructure and generates revenue from day one with traffic the platform already has; Layer 1 (direct platform inventory) once monthly platform traffic justifies a sales motion, typically 6–12 months post-launch; Layer 4 (creator tooling subscriptions) once power-user behavior patterns emerge from analytics, typically 6–9 months post-launch; and Layer 3 (sponsored discovery placements) last, only after the discovery surfaces have enough organic traffic that sponsored placements have value, typically 12+ months post-launch. Launching Layer 3 before discovery is mature creates a perverse incentive where the platform over-promotes sponsored content to manufacture the value proposition.

Chapter 6 • Earnings-Claim Compliance Language

The original concept used "GLOBAL EARNING PAGES" as the product name and "Earn Money" as the dashboard button label. Both create implied-earnings exposure under the U.S. Federal Trade Commission's endorsement rules (16 CFR Part 255) and analogous regimes in other jurisdictions. The FTC requires that earnings claims be substantiated by the same evidence that would be required if the advertiser made the claim directly, and that bloggers and endorsers are liable for misleading or unsubstantiated representations. The product name and UI labels are not exempt from these rules — the entire user-facing framing is examined for implied promises.

The compliance concern is not that the platform promises users a specific dollar amount. It is that the words "earn" and "earning" create an expectation of monetary return that the platform cannot guarantee, because earnings depend on factors entirely outside the platform's control: the user's ad-network approval, the user's traffic quality, the user's content quality, the ad network's payout rates, and the ad network's continued willingness to serve ads to the user's visitors. A platform that frames itself around "earning" while disclaiming any responsibility for actual earnings is in a structurally weak compliance position.

6.1 Recommended language changes

Original concept	Recommendation	Rationale
"GLOBAL EARNING PAGES"	"Global Creator Pages" or "Creator Hub"	Removes the earnings promise from the product name
"Earn Money" button	"Monetization" or "Ad Networks"	Neutral descriptor of the feature, not an outcome
"Earnings" dashboard section	"External Network Stats" with link-out to provider dashboard	Accurate: the platform is showing third-party data, not its own
"Your earnings" labels	"Your ad-network stats (provided by [Adsterra/Monetag])"	Makes clear the source and disclaims platform responsibility
Traffic-challenge rewards framed as "boost your earnings"	"Grow your audience" with rewards in XP/badges/themes	Reframes incentive around legitimate engagement, not ad-clicking
Creator Levels described as "earning tiers"	"Activity levels" with criteria based on engagement, not earnings	Avoids implying that higher level = higher income

Table 6.1 — Recommended language substitutions

6.2 Required disclaimers

The platform must publish, and require users to acknowledge, a clear earnings disclaimer before they activate any ad integration. The disclaimer must state, at minimum: that the platform does not pay users; that earnings depend entirely on the user's ad-network relationship; that the platform cannot guarantee any specific level of earnings or any earnings at all; that ad-network approval, traffic quality, content quality, and ad-network payout rates are outside the platform's control; and that any statistics displayed on the platform are sourced

from the ad network and are not the platform's own measurements of earnings. The disclaimer must be presented as a required step in the ad-integration onboarding flow, not buried in a terms-of-service document that users do not read.

The platform must also publish a separate content-policy statement that prohibits users from representing their page as a guaranteed earnings opportunity. Users who promote their Special Page externally with "earn money" framing, click-farm language, or "get paid to" messaging violate platform policy and must be subject to moderation action. This protects the platform from downstream liability for user misrepresentations and aligns with the ad networks' own prohibitions on incentivized traffic.

6.3 Analytics display rules

The platform must not display fabricated or estimated earnings numbers. If the ad network provides a verified API for earnings data, the platform may display that data with clear attribution to the source network. If the ad network does not provide a verified API, the platform must display only platform-side analytics (page views, visitors, traffic sources) and link the user to their ad-network dashboard for earnings information. The platform must not estimate earnings by multiplying page views by an assumed CPM, because such estimates can be construed as implied earnings claims and can diverge sharply from actual ad-network payouts.

Compliance principle: the platform measures traffic. The ad network measures earnings. The platform never reports earnings as its own data. When in doubt, show less monetary framing and link out.

Chapter 7 • Account Lifecycle and Offboarding

The original concept did not specify what happens when an ad-network account is suspended, when a user deletes their platform account, when a user wants to switch networks, or when the platform bans a user. Each of these is a common operational event that must be handled by the ad-integration lifecycle state machine. Failing to handle them produces dead ad code on live pages, which is at best a UX bug and at worst a compliance violation (e.g. serving ads for a banned Adsterra account).

Under the General Data Protection Regulation, individuals have a right to erasure of their personal data (Article 17) and the right to withdraw consent for processing. This right is not absolute — the platform may retain data necessary for legal compliance, fraud prevention, or establishment/defense of legal claims — but the default behavior on account deletion must be deletion of personal data, not silent retention. User-generated content (Special Pages, posts) is a particular challenge: retention is necessary to honor the user's audience that has bookmarked or shared the page, but personal-data elements (profile information, analytics fingerprints) must be removed or anonymized.

7.1 Ad-integration lifecycle states

Every user-supplied ad integration moves through a defined state machine. The state determines whether the integration's ad code is rendered on live pages, whether the integration can be edited, and what offboarding actions are required. States are: Draft (created but not submitted), PendingReview (submitted, awaiting platform review), Approved (live), Disabled (user-disabled or admin-disabled, but recoverable), Revoked (ad-network account suspended, integration not recoverable), and Deleted (user deleted the integration, code removed from all pages).

State	Ad code rendered?	User can edit?	Transition triggers
Draft	No	Yes	User submits for review
PendingReview	No	No (locked)	Admin approves or rejects
Approved	Yes (subject to placement engine)	Yes (re-review required)	User disables; admin disables; ad-network webhook reports suspension
Disabled	No	Yes (re-submit for review)	User re-enables; user deletes; admin bans
Revoked	No	No (integration is terminal)	User deletes; admin cleans up
Deleted	No (code purged from all pages)	No (record retained for audit)	Terminal state

Table 7.1 — Ad-integration lifecycle state machine

7.2 Offboarding scenarios

7.2.1 User's ad-network account is suspended

When an ad network reports (via webhook or admin notification) that a user's publisher account has been suspended, the platform immediately transitions all of that user's integrations with that network to Revoked state. The placement engine stops rendering the integration's ad code on all of the user's pages within 60 seconds (the placement engine cache TTL). The user is notified with the reason provided by the network, and is informed that the integration is not recoverable through the platform; they must resolve the suspension directly with the network. If the suspension is resolved at the network level and the network re-verifies the user's site, the user must re-submit the integration through the standard onboarding flow — the platform does not auto-revive Revoked integrations.

7.2.2 User deletes their platform account

Account deletion triggers a multi-step offboarding flow. First, all ad integrations are transitioned to Deleted state and their ad code is purged from all cached render outputs within 60 seconds. Second, all Special Pages are transitioned to a Tombstone state: the page URL returns a 410 Gone response for 30 days (to inform search engines and inbound links), then the page is fully removed. Third, personal data in the user profile is anonymized (name, email, avatar, biographical fields replaced with null or "deleted user"). Fourth, analytics events associated with the user are anonymized by replacing the userId with a hash that cannot be reversed. Fifth, financial records (if any subscription history exists) are retained for the period required by tax law, with personal-data fields stripped.

7.2.3 User wants to migrate from one network to another

Migration is handled by enabling both integrations simultaneously during the transition period, then disabling the old one. Because the placement engine supports multiple integrations on the same page (subject to compatibility rules), the new integration can be approved and activated before the old one is disabled, eliminating any gap in monetization. The user must explicitly disable the old integration; the platform does not auto-disable based on inactivity, because ad-network accounts in good standing may legitimately have periods of zero traffic.

7.2.4 Platform bans a user

Platform-initiated bans follow the same offboarding path as user-initiated deletion, with two exceptions. First, the user's pages are not tombstoned with a 410; they are immediately set to return a 403 Forbidden with a generic "page unavailable" message, to prevent the user from continuing to send traffic to a now-non-monetizable page. Second, the user's data is retained in full (not anonymized) for the period required by the platform's legal hold policy, typically 7 years, to support defense against any subsequent legal claim by the user. After the hold period, the data is anonymized under the same flow as user-initiated deletion.

7.3 Data retention policy

Data class	Default retention	On user deletion	Legal basis for retention
User profile data	Account lifetime	Anonymized within 30 days	GDPR Art. 17 — no longer necessary
Special Page content	Account lifetime	Tombstone 30 days, then hard delete	Honor inbound links; SEO cleanup

Data class	Default retention	On user deletion	Legal basis for retention
Ad-integration records	Account lifetime + 7 years	Anonymized except audit fields	Fraud prevention; ad-network audit obligations
Consent records	13 months (TCF v2.2 max)	Hard delete after expiry	TCF v2.2; demonstrate consent under GDPR Art. 7.1
Analytics events	26 months	Anonymized (userId → hash)	Aggregate trend analysis; no individual profiling after deletion
Moderation events	Account lifetime + 7 years	Retained in full	Defense against legal claims; demonstrate moderation decisions
Financial records (subscriptions)	7 years post-transaction	Retained in full	Tax law; financial record-keeping requirements

Table 7.2 — Data retention policy by class

Chapter 8 • Consolidated Data Model Changes

The preceding chapters introduce a number of new tables and columns. This chapter consolidates them into a single reference, organized by the issue they address. The intent is to give the engineering team a complete migration plan: every new table, every new column, every new enum value, in one place, with the issue that drove the change noted alongside.

8.1 Issue 1 — Co-display rules

New tables: AdNetwork, AdNetworkCompatibility, AdPlacementPolicy. New columns on existing tables: AdPlacement gains priority (Int) and contentClassRequired (String?) to support the placement engine's compatibility checks. New enum: CompatibilityVerdict (ALLOWED, ALLOWED_WITH_LIMITS, FORBIDDEN), AdNetworkStatus (ACTIVE, DEPRECATED, BANNED).

8.2 Issue 2 — Consent layer

New tables: ConsentRecord, AdNetworkVendorMapping. New columns on existing tables: Visitor gains jurisdiction (String, derived from GeoIP) and consentState (FK to most recent ConsentRecord). The SpecialPage table gains requiresConsent (Bool, default true) and cmpConfigOverride (Json?, for per-page CMP customization, normally null).

8.3 Issue 3 — Trust Score

New tables: TrustScore (snapshot per page, written nightly), ModerationEvent (append-only state-transition log), SignalSource (registry of signal providers and their config). The SpecialPage table gains moderationState (enum: Pending, Approved, Restricted, Suspended, Banned) replacing the previous simple Boolean isApproved field.

8.4 Issue 4 — Revenue model

New tables: PlatformAdCampaign (for Layer 1 direct-sold inventory), PlatformAdNetworkIntegration (for Layer 2 platform's own ad-network account), SponsoredPlacement (for Layer 3 discovery sponsorship), CreatorSubscription (for Layer 4 creator tooling). The AdPlacement table gains source (enum: USER_INTEGRATION, PLATFORM_DIRECT, PLATFORM_NETWORK) so the placement engine knows which pool of inventory a given placement draws from.

8.5 Issue 5 — Compliance language

No structural data model changes. UI string changes are managed through the i18n catalog. One new table is required: EarningsDisclaimerAck — a simple append-only log recording which users have acknowledged the earnings disclaimer, with timestamp and disclaimer version, to support audit defense.

8.6 Issue 6 — Lifecycle & offboarding

The AdIntegration table (already present in the original concept) gains lifecycleState (enum: Draft, PendingReview, Approved, Disabled, Revoked, Deleted) replacing the previous simple status field. New

tables: OffboardingJob (tracks multi-step offboarding flows, including the 30-day tombstone timer for deleted accounts), LegalHold (records the legal basis and expiry for retained data on banned users).

Chapter 9 • Updated Rollout Sequence

The original concept did not sequence the launch. This chapter proposes a three-phase rollout with explicit risk gates between phases. The objective is to launch with the minimum feature set that demonstrates the platform thesis, then add revenue layers and advanced moderation as traffic and team capacity justify. Each phase has a defined exit criterion; phases do not advance until exit criteria are met.

9.1 Phase 1 — Minimum viable platform (MVP)

Phase 1 delivers the minimum feature set required to validate the thesis: users can create Special Pages, publish content, share URLs, and connect one supported ad network (Adsterra). Platform ads run via Layer 2 (platform's own Adsterra account) on eligible pages. The Consent Layer is live with a certified CMP. The Trust Score runs in shadow mode (computed but not acted on) for the first 60 days to allow calibration. Manual approval is required for every ad integration. Moderation is fully manual.

Capability	Phase 1 status	Exit criterion
Special Page creation	Live	1,000 published pages
Page builder (basic blocks)	Live	NPS > 30 from page creators
Adsterra integration	Live (manual approval)	50 approved integrations
Monetag integration	Deferred to Phase 2	—
Platform ads (Layer 2)	Live	Platform CPM stable for 30 days
Consent Layer (CMP)	Live (IAB TCF v2.2)	0 consent-related compliance incidents
Trust Score	Shadow mode	Calibration against manual review > 80% agreement
Moderation	Fully manual	24-hour median response time
i18n	English only	—
Christmas 2026 campaign	Live as flagship campaign	—

Table 9.1 — Phase 1 capabilities and exit criteria

9.2 Phase 2 — Multi-network and moderation automation

Phase 2 adds Monetag as a second supported network, with the placement engine's compatibility rules fully live. The Trust Score transitions from shadow mode to active: scores below threshold automatically transition pages to Restricted state, with moderator review of every automated transition for the first 90 days. Layer 4 (creator tooling subscriptions) launches with a basic tier. Internationalization expands to French, Spanish, and Portuguese.

9.3 Phase 3 — Direct sales and discovery

Phase 3 launches Layer 1 (direct platform inventory sold by a platform sales team) and Layer 3 (sponsored discovery placements). The two are launched together because Layer 1 inventory becomes more valuable when sponsored discovery drives targeted traffic to specific pages, and Layer 3 sponsored placements become more valuable when Layer 1 ad inventory demonstrates the platform's ability to convert traffic. By this point the Trust Score has been calibrated for 12+ months and automated moderation handles the bulk of routine decisions, with human moderators focused on appeals and edge cases. Internationalization adds Arabic, Hindi, Swahili, and German.

Chapter 10 • Remaining Open Items

Resolving the six issues raised in the original review does not close every open question. The following items are flagged for future work but are not blocking for Phase 1 launch. They should be revisited as the platform matures and as operational experience reveals which of them are real constraints and which are theoretical concerns.

10.1 Ad-network API integration depth

This document assumes the platform displays external ad-network stats with link-out to the ad-network dashboard, rather than integrating verified earnings APIs. If verified APIs become available and reliable, the analytics display rules in Chapter 6 should be revisited to allow tighter integration. This is a Phase 2+ consideration; Phase 1 should ship with link-out only.

10.2 Additional ad networks

The architecture is designed to support additional ad networks beyond Adsterra and Monetag, but each new network requires (a) an entry in the AdNetwork table, (b) an IAB TCF vendor ID mapping, (c) compatibility rules against every existing network, and (d) integration-type support. The order in which networks are added should be driven by user demand rather than platform preference. Google AdSense is the obvious next candidate but its compliance requirements are stricter than Adsterra or Monetag and may necessitate additional placement-engine work.

10.3 Mobile app vs mobile web

The original concept did not specify whether Special Pages are consumed primarily via mobile web or via a platform mobile app. The architecture assumes mobile web first, which simplifies the ad-network integration (ad tags work in web views; they do not work natively in mobile apps without an SDK). If a mobile app is added later, the in-app ad experience will require either an SDK integration per network or a webview-based fallback — neither is trivial and the decision should be deferred until mobile-app traffic justifies the investment.

10.4 Tax and reporting obligations for subscriptions

Layer 4 (creator tooling subscriptions) introduces tax-reporting obligations that do not exist for the ad-supported layers. Depending on the jurisdictions of paying users, the platform may need to collect VAT, GST, or sales tax, and may need to file periodic returns in multiple jurisdictions. This is a Phase 2 concern (when Layer 4 launches) but should be scoped with a tax advisor before Phase 2 begins.

10.5 Content licensing and IP

The original concept did not address how the platform handles content that infringes third-party intellectual property. User-uploaded images, videos, and music in Special Pages are a copyright liability if the platform does not have a DMCA-compliant takedown process. This is necessary infrastructure for any UGC platform operating in jurisdictions with safe-harbor provisions, and should be scoped before Phase 1 launch rather than added reactively after the first takedown request.